

Wei Zhang

Ann Arbor, MI | (734) 555-0164 | wei.zhang@example.com | github.com/weizhang-ml | weizhang-ml.github.io

Education

Sep 2022 – Expected May 2027 Ann Arbor, MI	Ph.D. in Computer Science and Engineering <i>University of Michigan</i>
Sep 2018 – Jun 2022 Beijing, China	B.Eng. in Automation <i>Tsinghua University</i> GPA: 3.8/4.0

Research Experience

Sep 2022 – Present Ann Arbor, MI	Graduate Research Assistant <i>Michigan AI Lab</i> - Study how retrieval-augmented language models fail on long documents, advised by Prof. Maria Torres - Built an open benchmark of 12,000 questions over book-length texts, now used by 4 other labs - Designed a sparse mixture-of-experts retriever that cuts memory use by 60% at the same accuracy
May 2025 – Aug 2025 Redmond, WA	Research Intern <i>Microsoft Research</i> Added conformal prediction to a production question answering model so it abstains when unsure, cutting wrong answers by 35%
Jan 2024 – May 2024 Ann Arbor, MI	Graduate Student Instructor, Introduction to Machine Learning <i>University of Michigan</i> Taught two weekly discussion sections of 40 students and wrote 3 new homework problems on decision trees and kernels

Publications

2026	W. Zhang, R. Patel, and M. Torres, “Sparse Mixture-of-Experts for Long-Context Retrieval,” <i>International Conference on Learning Representations (ICLR)</i> , 2026. arxiv.org/abs/2510.01234
Dec 2025	W. Zhang and M. Torres, “Calibrating Retrieval-Augmented Models with Conformal Prediction,” <i>Advances in Neural Information Processing Systems (NeurIPS)</i> , pp. 1123–1135, Dec 2025.
Jul 2025	J. Kim, W. Zhang, A. Gupta, and M. Torres, “Benchmarking Long-Document Question Answering,” <i>Proc. Annual Meeting of the Association for Computational Linguistics (ACL)</i> , Vienna, Austria, pp. 410–422, Jul 2025. doi: 10.18653/v1/2025.acl-long.33.

Honors & Awards

2023	NSF Graduate Research Fellowship , National Science Foundation
Dec 2024	Outstanding Reviewer Award , NeurIPS

Skills

Machine Learning	PyTorch, JAX, Hugging Face Transformers, retrieval-augmented generation
Programming	Python, C++, CUDA, SQL
Tools	Slurm, Docker, Weights & Biases, LaTeX