

Wei Zhang

Ann Arbor, MI | (734) 555-0164 | wei.zhang@example.com | github.com/weizhang-ml | weizhang-ml.github.io

EDUCATION

University of Michigan

Ph.D. in Computer Science and Engineering

Ann Arbor, MI

Sep 2022 – Expected May 2027

Tsinghua University

B.Eng. in Automation, GPA: 3.8/4.0

Beijing, China

Sep 2018 – Jun 2022

RESEARCH EXPERIENCE

Graduate Research Assistant

Michigan AI Lab

Sep 2022 – Present

Ann Arbor, MI

- Study how retrieval-augmented language models fail on long documents, advised by Prof. Maria Torres
- Built an open benchmark of 12,000 questions over book-length texts, now used by 4 other labs
- Designed a sparse mixture-of-experts retriever that cuts memory use by 60% at the same accuracy

Research Intern

Microsoft Research

May 2025 – Aug 2025

Redmond, WA

- Added conformal prediction to a production question answering model so it abstains when unsure, cutting wrong answers by 35%

Graduate Student Instructor, Introduction to Machine Learning

University of Michigan

Jan 2024 – May 2024

Ann Arbor, MI

- Taught two weekly discussion sections of 40 students and wrote 3 new homework problems on decision trees and kernels

PUBLICATIONS

- [1] W. Zhang, R. Patel, and M. Torres, “Sparse Mixture-of-Experts for Long-Context Retrieval,” *International Conference on Learning Representations (ICLR)*, 2026. arxiv.org/abs/2510.01234
- [2] W. Zhang and M. Torres, “Calibrating Retrieval-Augmented Models with Conformal Prediction,” *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 1123–1135, Dec 2025.
- [3] J. Kim, W. Zhang, A. Gupta, and M. Torres, “Benchmarking Long-Document Question Answering,” *Proc. Annual Meeting of the Association for Computational Linguistics (ACL)*, Vienna, Austria, pp. 410–422, Jul 2025. doi: 10.18653/v1/2025.acl-long.33.

HONORS & AWARDS

NSF Graduate Research Fellowship, National Science Foundation

2023

Outstanding Reviewer Award, NeurIPS

Dec 2024

TECHNICAL SKILLS

Machine Learning: PyTorch, JAX, Hugging Face Transformers, retrieval-augmented generation

Programming: Python, C++, CUDA, SQL

Tools: Slurm, Docker, Weights & Biases, LaTeX