

Wei Zhang

Ann Arbor, MI | (734) 555-0164 | wei.zhang@example.com | github.com/weizhang-ml | weizhang-ml.github.io

EDUCATION

University of Michigan <i>Ph.D. in Computer Science and Engineering</i>	Sep 2022 – Expected May 2027 Ann Arbor, MI
Tsinghua University <i>B.Eng. in Automation, GPA: 3.8/4.0</i>	Sep 2018 – Jun 2022 Beijing, China

RESEARCH EXPERIENCE

Graduate Research Assistant <i>Michigan AI Lab</i> - Study how retrieval-augmented language models fail on long documents, advised by Prof. Maria Torres - Built an open benchmark of 12,000 questions over book-length texts, now used by 4 other labs - Designed a sparse mixture-of-experts retriever that cuts memory use by 60% at the same accuracy	Sep 2022 – Present Ann Arbor, MI
Research Intern <i>Microsoft Research</i> Added conformal prediction to a production question answering model so it abstains when unsure, cutting wrong answers by 35%	May 2025 – Aug 2025 Redmond, WA
Graduate Student Instructor, Introduction to Machine Learning <i>University of Michigan</i> Taught two weekly discussion sections of 40 students and wrote 3 new homework problems on decision trees and kernels	Jan 2024 – May 2024 Ann Arbor, MI

PUBLICATIONS

[1] W. Zhang, R. Patel, and M. Torres, “Sparse Mixture-of-Experts for Long-Context Retrieval,” *International Conference on Learning Representations (ICLR)*, 2026. arxiv.org/abs/2510.01234

[2] W. Zhang and M. Torres, “Calibrating Retrieval-Augmented Models with Conformal Prediction,” *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 1123–1135, Dec 2025.

[3] J. Kim, W. Zhang, A. Gupta, and M. Torres, “Benchmarking Long-Document Question Answering,” *Proc. Annual Meeting of the Association for Computational Linguistics (ACL)*, Vienna, Austria, pp. 410–422, Jul 2025. doi: 10.18653/v1/2025.acl-long.33.

HONORS & AWARDS

NSF Graduate Research Fellowship , National Science Foundation	2023
Outstanding Reviewer Award , NeurIPS	Dec 2024

SKILLS

Machine Learning: PyTorch, JAX, Hugging Face Transformers, retrieval-augmented generation

Programming: Python, C++, CUDA, SQL

Tools: Slurm, Docker, Weights & Biases, LaTeX